



The New Java™ Technology Memory Model

java.sun.com/javaone/sf

Jeremy Manson and William Pugh
<http://www.cs.umd.edu/~pugh>



Audience

- Assume you are familiar with basics of Java™ technology-based threads (“Java threads”)
 - Creating, starting and joining threads
 - Synchronization
 - `wait` and `notifyAll`

Java Thread Specification

- Revised as part of JSR-133
- Part of the new Java Language Spec
 - and the Virtual Machine Spec
- Features talked about here today are in JDK1.5
 - Not all of these ideas are guaranteed to work in previous versions
 - Previous thread spec was broken
 - forbid optimizations performed by many JVMs

Safety Issues in Multithreaded Systems

- Many intuitive assumptions do not hold
- Some widely used idioms are not safe
 - Original Double-checked locking idiom
 - Checking non-volatile flag for thread termination
- Can't use testing to check for errors
 - Some anomalies will occur only on some platforms
 - e.g., multiprocessors
 - Anomalies will occur rarely and non-repeatedly

Revising the Thread Spec

- The Java Thread Specification has undergone significant revision
 - Mostly to correctly formalize existing behavior
 - But a few changes in behavior
- Goals
 - Clear and easy to understand
 - Foster reliable multithreaded code
 - Allow for high performance JVMs
- Has affected JVMs
 - And badly written existing code
 - Including parts of Sun's JDK

This Talk...

- Describe building blocks of synchronization and concurrent programming in Java
 - Both language primitives and `util.concurrent` abstractions
- Explain what it means for code to be correctly synchronized
- Try to convince you that clever reasoning about unsynchronized code is almost certainly wrong
 - Not needed for efficient and reliable programs

This Talk...

- We will be talking mostly about
 - synchronized methods and blocks
 - volatile fields
- Same principles work with JSR-166 locks and atomic operations
- Will also talk about final fields and immutability.

Taxonomy

- High level concurrency abstractions
 - JSR-166 and `java.util.concurrent`
- Low level locking
 - `synchronized()` blocks
- Low level primitives
 - volatile variables, `java.util.concurrent.atomic` classes
 - allows for non-blocking synchronization
- Data races: deliberate undersynchronization
 - Avoid!
 - Not even Doug Lea can get it right


Three Aspects of Synchronization

- Atomicity
 - Locking to obtain mutual exclusion
- Visibility
 - Ensuring that changes to object fields made in one thread are seen in other threads
- Ordering
 - Ensuring that you aren't surprised by the order in which statements are executed

Don't Try To Be Too Clever


- People worry about the cost of synchronization
 - Try to devise schemes to communicate between threads without using synchronization
 - locks, volatiles, or other concurrency abstractions
- Nearly impossible to do correctly
 - Inter-thread communication without synchronization is not intuitive

Quiz Time



Can this result in $i = 0$ and $j = 0$?

Answer: Yes!



How can $i = 0$ and $j = 0$?

How Can This Happen?

- Compiler can reorder statements
 - Or keep values in registers
- Processor can reorder them
- On multi-processor, values not synchronized in global memory
- The memory model is designed to allow aggressive optimization
 - including optimizations no one has implemented yet
- Good for performance
 - bad for your intuition about insufficiently synchronized code

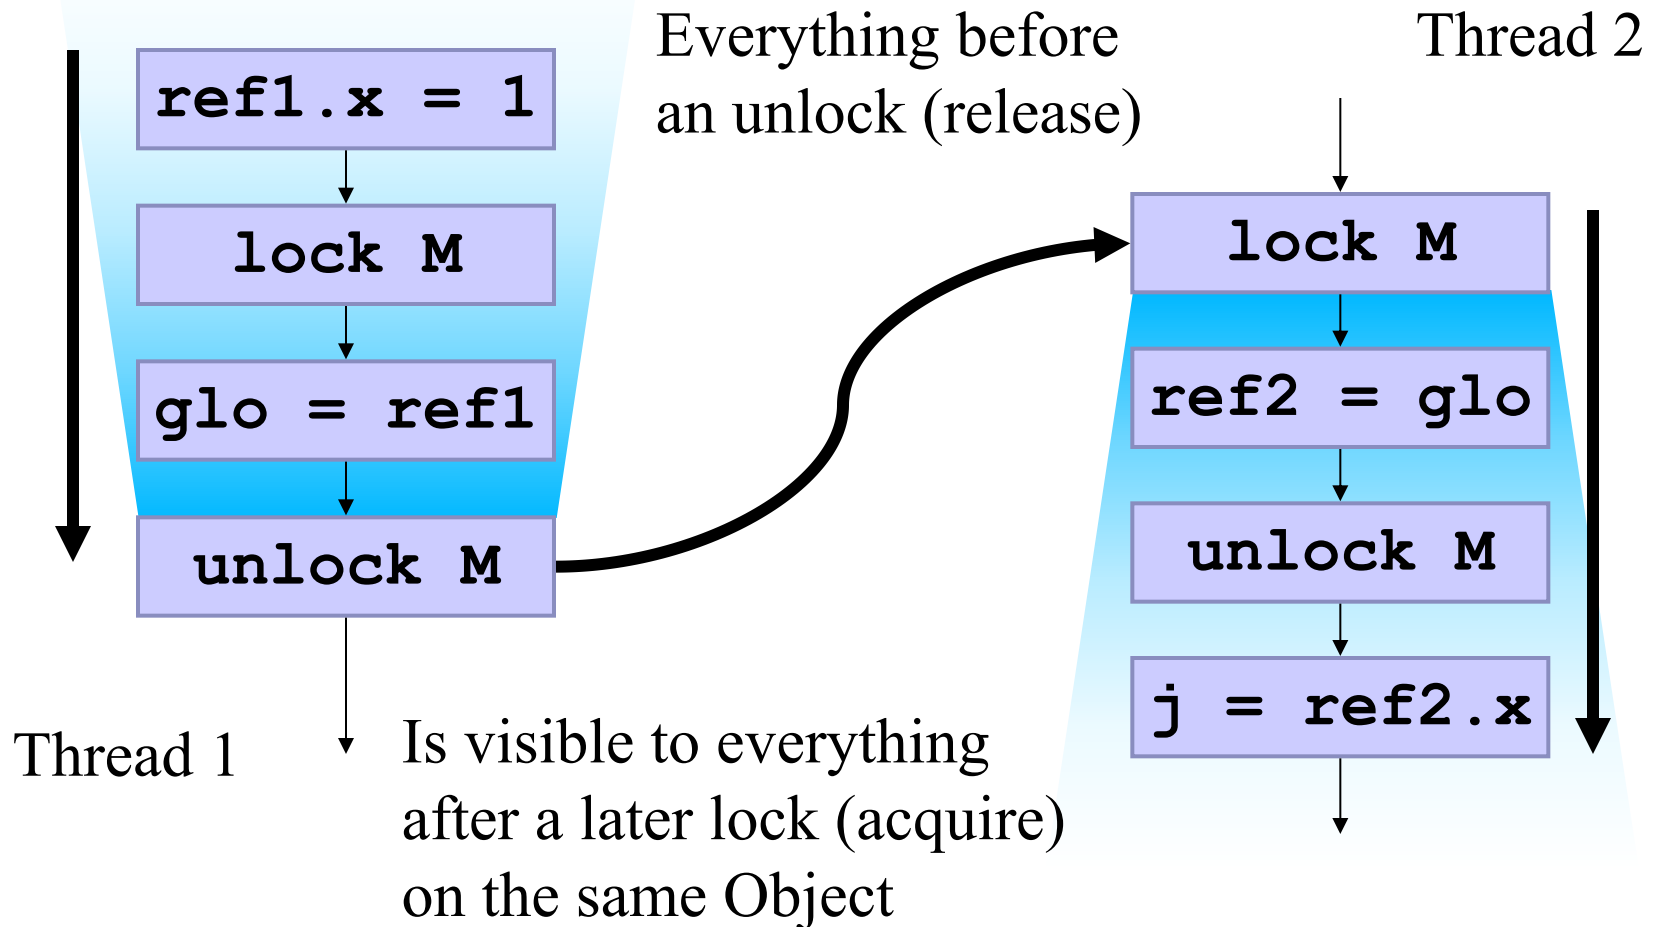
Correctness and Optimizations

- Clever code that depends the order you think the system *must* do things in is almost always wrong in Java
- Dekker's Algorithm (first correct lock implementation) requires this ordering
 - doesn't work in Java, use supplied locks
- Must use synchronization to enforce visibility and ordering
 - As well as mutual exclusion
 - If you use synchronization correctly, you will not be able to see reorderings

Synchronization Actions (*approximately*)

```
// block until obtain lock  
synchronized(anObject) {  
    // get main memory value of field1 and field2  
    int x = anObject.field1;  
    int y = anObject.field2;  
    anObject.field3 = x+y;  
    // commit value of field3 to main memory  
}  
  
// release lock  
moreCode();
```

When Are Actions Visible to Other Threads?



Release and Acquire

- All accesses before a release
 - are ordered before and visible to
 - any accesses after a matching acquire
- Unlocking a monitor/lock is a release
 - that is acquired by any following lock of *that* monitor/lock

Ordering

- Roach motel ordering
 - Compiler/processor can move accesses into synchronized blocks
 - Can only move them out under special circumstances, generally not observable
- Some special cases:
 - locks on thread local objects are a no-op
 - reentrant locks are a no-op

Volatile fields

- If a field could be simultaneously accessed by multiple threads, and at least one of those accesses is a write
 - make the field volatile
 - documentation
 - gives essential JVM guarantees
 - Can be tricky to get right, but nearly impossible without volatile
- What does volatile do?
 - reads and writes go directly to memory
 - not cached in registers
 - volatile longs and doubles are atomic
 - not true for non-volatile longs and doubles
 - compiler reordering of volatile accesses is restricted
 - roach motel semantics for volatiles and normals
 - no reordering for volatiles and volatiles

Volatile release/acquire

- A volatile write is a release
 - that is acquired by a later read of the same variable
- All accesses before the volatile write
 - are ordered before and visible to all accesses after the volatile read

Volatile guarantees visibility

- **stop** *must* be declared volatile
 - Otherwise, compiler could keep in register

```
class Animator implements Runnable {  
    private volatile boolean stop = false;  
    public void stop() { stop = true; }  
    public void run() {  
        while (!stop)  
            oneStep();  
    }  
    private void oneStep() { /*...*/ }  
}
```

Volatile guarantees ordering

- If a thread reads **data**, there is a release/acquire on **ready** that guarantees visibility and ordering

```
class Future {  
    private volatile boolean ready;  
    private Object data;  
    public Object get() {  
        if (!ready)  
            return null;  
        return data;  
    }  
  
    public synchronized  
        void setOnce(Object o) {  
        if (ready) throw ... ;  
        data = o;  
        ready = true;  
    }  
}
```

Other Acquires and Releases

- Other actions form release/acquire pairs
- Starting a thread is a release
 - acquired by the run method of the thread
- Termination of a thread is a release
 - acquired by any thread that joins with the terminated thread

Defending against data races

- Attackers can pass instances of your object to other threads via a data race
- Can cause weird things to be observed
 - could be observed in some JVMs
 - in older JVMs, `String` objects might be seen to change
 - change from `/tmp` to `/usr`
- If a class is security critical, must take steps
- Choices:
 - use synchronization (even in constructor)
 - object can be made visible to multiple threads before constructor finishes
 - make object immutable by making all fields final

Immutable classes

- Make all critical fields final
- Don't allow other threads to see object until it is fully constructed
- JVM will be responsible for ensuring that object is perceived as immutable
 - even if malicious code uses data races to attack the class

Optimization of final fields

- New spec allows aggressive optimization of final fields
 - hoisting of reads of final fields across synchronization and unknown method calls
 - still maintains immutability
- Should allow for future JVMs to obtain performance advantages

Finalizers

- Only guaranteed to see writes that occur by the end of the object's constructor.
 - If finalizer needs to see later writes, use synchronization
- Fields may be made final earlier than the program text might imply
 - Synchronization on object also keeps it alive
- Multiple finalizers may be run concurrently
 - Be careful to synchronize properly!

Synchronize When Needed

- Places where threads interact
 - Need synchronization
 - May need careful thought
 - May need documentation
 - Cost of required synchronization not significant
 - For most applications
 - No need to get tricky

Synchronized Classes

- Some classes are synchronized
 - **Vector**, **Hashtable**, **Stack**
 - Most Input/Output Streams
 - Overhead of unneeded synchronization can be measurable
- Contrast with Collection classes
 - By default, not synchronized
 - Can request synchronized version
 - Or can use `java.util.concurrent` versions (**Queue**, **ConcurrentMap** implementations)
- Using synchronized classes
 - Often doesn't suffice for concurrent interaction

Synchronized Collections Aren't Always Enough

- Transactions (DO NOT USE)
 - Violate atomicity...

```
ID getID(String name) {  
    ID x = h.get(name);  
    if (x == null) {  
        x = new ID();  
        h.put(name, x);  
    }  
    return x;  
}
```

- Iterators
 - Can't modify collection while another thread is iterating through it

Concurrent Interactions

- Often need entire transactions to be atomic
 - Reading and updating a Map
 - Writing a record to an OutputStream
- OutputStreams are synchronized
 - Can have multiple threads trying to write to the same OutputStream
 - Output from each thread is nondeterministically interleaved
 - Often essentially useless

util.concurrent

- The stuff in `java.util.concurrent` is great, use it
- **ConcurrentHashMap** has some additional features to get around problems with transactions
 - `putIfAbsent`
 - concurrent iteration
- **CopyOnWrite** classes allow concurrent iteration and non-blocking reads
 - modification is expensive, should be rare

Designing Fast Code

- Make it right before you make it fast
- Reduce synchronization costs
 - Avoid sharing mutable objects across threads
 - avoid old Collection classes (**Vector**, **Hashtable**)
 - use bulk I/O (or, even better, java.nio classes)
- Use java.util.concurrent classes
 - designed for speed, scalability and correctness
- Avoid lock contention
 - Reduce lock scopes
 - Reduce lock durations

Things That Don't Work

- Thinking about memory barriers
 - There is nothing that gives you the effect of a memory barrier
- Original Double-Check Idiom
 - AKA multithreaded lazy initialization
 - Any unsynchronized non-volatile reads/writes of refs
- Depending on sleep for visibility
- Clever reasoning about cause and effect with respect to data races

Synchronization on Thread Local Objects

- Synchronization on thread local objects
 - (objects that are only accessed by a single thread)
 - has no semantics or meaning
 - compiler can remove it
 - can also remove reentrant synchronization
 - e.g., calling a synchronized method from another synchronized method on same object
- This is an optimization people have talked about for a while
 - not sure if anyone is doing it yet

Thread safe lazy initialization

- Want to perform lazy initialization of something that will be shared by many threads
- Don't want to pay for synchronization after object is initialized
- Standard double-checked locking doesn't work
 - making the checked field volatile fixes it
- If two threads might simultaneously access a field, and one of them writes to it
 - the field must be volatile

Wrap-up

- Cost of synchronization operations can be significant
 - But cost of *needed* synchronization rarely is
- Thread interaction needs careful thought
 - But not too clever
 - Don't want to have to think too hard about reordering
 - No data races in your program, no observable reordering
- Need for inter-thread communication...

Wrap-up - Communication

- Communication between threads
 - Requires *both* threads to interact via synchronization
- JSR-133 & 166 provide new mechanisms for communication
 - High level concurrency framework
 - volatile fields
 - final fields

Q&A

